

10.6 Computer Graphics

Computer graphics deals with images. The images are moved around. Their scale is changed. Three dimensions are projected onto two dimensions. All the main operations are done by matrices—but the shape of these matrices is surprising.

The transformations of three-dimensional space are done with 4 by 4 matrices. You would expect 3 by 3. The reason for the change is that one of the four key operations cannot be done with a 3 by 3 matrix multiplication. Here are the four operations:

- Translation** (shift the origin to another point $P_0 = (x_0, y_0, z_0)$)
- Rescaling** (by c in all directions or by different factors c_1, c_2, c_3)
- Rotation** (around an axis through the origin or an axis through P_0)
- Projection** (onto a plane through the origin or a plane through P_0).

Translation is the easiest—just add (x_0, y_0, z_0) to every point. But this is not linear! No 3 by 3 matrix can move the origin. So we change the coordinates of the origin to $(0, 0, 0, 1)$. This is why the matrices are 4 by 4. The “homogeneous coordinates” of the point (x, y, z) are $(x, y, z, 1)$ and we now show how they work.

1. Translation Shift the whole three-dimensional space along the vector \mathbf{v}_0 . The origin moves to (x_0, y_0, z_0) . This vector \mathbf{v}_0 is added to every point \mathbf{v} in \mathbf{R}^3 . Using homogeneous coordinates, the 4 by 4 matrix T shifts the whole space by \mathbf{v}_0 :

$$\text{Translation matrix} \quad T = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ x_0 & y_0 & z_0 & 1 \end{bmatrix}.$$

Important: *Computer graphics works with row vectors.* We have row times matrix instead of matrix times column. You can quickly check that $[0 \ 0 \ 0 \ 1]T = [x_0 \ y_0 \ z_0 \ 1]$.

To move the points $(0, 0, 0)$ and (x, y, z) by \mathbf{v}_0 , change to homogeneous coordinates $(0, 0, 0, 1)$ and $(x, y, z, 1)$. Then multiply by T . A row vector times T gives a row vector. **Every \mathbf{v} moves to $\mathbf{v} + \mathbf{v}_0$:** $[x \ y \ z \ 1]T = [x + x_0 \ y + y_0 \ z + z_0 \ 1]$.

The output tells where any \mathbf{v} will move. (It goes to $\mathbf{v} + \mathbf{v}_0$.) Translation is now achieved by a matrix, which was impossible in \mathbf{R}^3 .

2. Scaling To make a picture fit a page, we change its width and height. A copier will rescale a figure by 90%. In linear algebra, we multiply by .9 times the identity matrix. That matrix is normally 2 by 2 for a plane and 3 by 3 for a solid. In computer graphics, with homogeneous coordinates, the matrix is *one size larger*:

$$\text{Rescale the plane:} \quad S = \begin{bmatrix} .9 & & \\ & .9 & \\ & & 1 \end{bmatrix}$$

$$\text{Rescale a solid:} \quad S = \begin{bmatrix} c & 0 & 0 & 0 \\ 0 & c & 0 & 0 \\ 0 & 0 & c & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}.$$

Important: S is not cI . We keep the “1” in the lower corner. Then $[x, y, 1]$ times S is the correct answer in homogeneous coordinates. The origin stays in its normal position because $[0\ 0\ 1]S = [0\ 0\ 1]$.

If we change that 1 to c , the result is strange. **The point (cx, cy, cz, c) is the same as $(x, y, z, 1)$.** The special property of homogeneous coordinates is that *multiplying by cI does not move the point*. The origin in \mathbf{R}^3 has homogeneous coordinates $(0, 0, 0, 1)$ and $(0, 0, 0, c)$ for every nonzero c . This is the idea behind the word “homogeneous.”

Scaling can be different in different directions. To fit a full-page picture onto a half-page, scale the y direction by $\frac{1}{2}$. To create a margin, scale the x direction by $\frac{3}{4}$. The graphics matrix is diagonal but not 2 by 2. It is 3 by 3 to rescale a plane and 4 by 4 to rescale a space:

$$\text{Scaling matrices} \quad S = \begin{bmatrix} \frac{3}{4} & & \\ & \frac{1}{2} & \\ & & 1 \end{bmatrix} \quad \text{and} \quad S = \begin{bmatrix} c_1 & & \\ & c_2 & \\ & & c_3 \\ & & & 1 \end{bmatrix}.$$

That last matrix S rescales the x, y, z directions by positive numbers c_1, c_2, c_3 . The extra column in all these matrices leaves the extra 1 at the end of every vector.

Summary The scaling matrix S is the same size as the translation matrix T . They can be multiplied. To translate and then rescale, multiply vTS . To rescale and then translate, multiply vST . Are those different? *Yes*.

The point (x, y, z) in \mathbf{R}^3 has homogeneous coordinates $(x, y, z, 1)$ in \mathbf{P}^3 . This “projective space” is not the same as \mathbf{R}^4 . It is still three-dimensional. To achieve such a thing, (cx, cy, cz, c) is the same point as $(x, y, z, 1)$. Those points of projective space \mathbf{P}^3 are really lines through the origin in \mathbf{R}^4 .

Computer graphics uses **affine** transformations, *linear plus shift*. An affine transformation T is executed on \mathbf{P}^3 by a 4 by 4 matrix with a special fourth column:

$$A = \begin{bmatrix} a_{11} & a_{12} & a_{13} & 0 \\ a_{21} & a_{22} & a_{23} & 0 \\ a_{31} & a_{32} & a_{33} & 0 \\ a_{41} & a_{42} & a_{43} & 1 \end{bmatrix} = \begin{bmatrix} T(1, 0, 0) & 0 \\ T(0, 1, 0) & 0 \\ T(0, 0, 1) & 0 \\ T(0, 0, 0) & 1 \end{bmatrix}.$$

The usual 3 by 3 matrix tells us three outputs, this tells four. The usual outputs come from the inputs $(1, 0, 0)$ and $(0, 1, 0)$ and $(0, 0, 1)$. When the transformation is linear, three outputs reveal everything. When the transformation is affine, the matrix also contains the output from $(0, 0, 0)$. Then we know the shift.

3. Rotation A rotation in \mathbf{R}^2 or \mathbf{R}^3 is achieved by an orthogonal matrix Q . The determinant is $+1$. (With determinant -1 we get an extra reflection through a mirror.) Include the extra column when you use homogeneous coordinates!

$$\text{Plane rotation} \quad Q = \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix} \quad \text{becomes} \quad R = \begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix}.$$

This matrix rotates the plane around the origin. *How would we rotate around a different point* $(4, 5)$? The answer brings out the beauty of homogeneous coordinates. *Translate* $(4, 5)$ *to* $(0, 0)$, *then rotate by* θ , *then translate* $(0, 0)$ *back to* $(4, 5)$:

$$vT_-RT_+ = \begin{bmatrix} x & y & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ -4 & -5 & 1 \end{bmatrix} \begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 4 & 5 & 1 \end{bmatrix}.$$

I won't multiply. The point is to apply the matrices one at a time: v translates to vT_- , then rotates to vT_-R , and translates back to vT_-RT_+ . Because each point $\begin{bmatrix} x & y & 1 \end{bmatrix}$ is a row vector, T_- acts first. The center of rotation $(4, 5)$ —otherwise known as $(4, 5, 1)$ —moves first to $(0, 0, 1)$. Rotation doesn't change it. Then T_+ moves it back to $(4, 5, 1)$. All as it should be. The point $(4, 6, 1)$ moves to $(0, 1, 1)$, then turns by θ and moves back.

In three dimensions, every rotation Q turns around an axis. The axis doesn't move—it is a line of eigenvectors with $\lambda = 1$. Suppose the axis is in the z direction. The 1 in Q is to leave the z axis alone, the extra 1 in R is to leave the origin alone:

$$Q = \begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \quad \text{and} \quad R = \begin{bmatrix} & & 0 \\ & Q & \\ 0 & 0 & 0 & 1 \end{bmatrix}.$$

Now suppose the rotation is around the unit vector $\mathbf{a} = (a_1, a_2, a_3)$. With this axis \mathbf{a} , the rotation matrix Q which fits into R has three parts:

$$Q = (\cos \theta)I + (1 - \cos \theta) \begin{bmatrix} a_1^2 & a_1a_2 & a_1a_3 \\ a_1a_2 & a_2^2 & a_2a_3 \\ a_1a_3 & a_2a_3 & a_3^2 \end{bmatrix} - \sin \theta \begin{bmatrix} 0 & a_3 & -a_2 \\ -a_3 & 0 & a_1 \\ a_2 & -a_1 & 0 \end{bmatrix}. \quad (1)$$

The axis doesn't move because $\mathbf{a}Q = \mathbf{a}$. When $\mathbf{a} = (0, 0, 1)$ is in the z direction, this Q becomes the previous Q —for rotation around the z axis.

The linear transformation Q always goes in the upper left block of R . Below it we see zeros, because rotation leaves the origin in place. When those are not zeros, the transformation is affine and the origin moves.

4. Projection In a linear algebra course, most planes go through the origin. In real life, most don't. A plane through the origin is a vector space. The other planes are affine spaces, sometimes called “flats.” An affine space is what comes from translating a vector space.

We want to project three-dimensional vectors onto planes. Start with a plane through the origin, whose unit normal vector is \mathbf{n} . (We will keep \mathbf{n} as a column vector.) The vectors in the plane satisfy $\mathbf{n}^T \mathbf{v} = 0$. *The usual projection onto the plane is the matrix* $I - \mathbf{n}\mathbf{n}^T$. To project a vector, multiply by this matrix. The vector \mathbf{n} is projected to zero, and the in-plane vectors \mathbf{v} are projected onto themselves:

$$(I - \mathbf{n}\mathbf{n}^T)\mathbf{n} = \mathbf{n} - \mathbf{n}(\mathbf{n}^T \mathbf{n}) = \mathbf{0} \quad \text{and} \quad (I - \mathbf{n}\mathbf{n}^T)\mathbf{v} = \mathbf{v} - \mathbf{n}(\mathbf{n}^T \mathbf{v}) = \mathbf{v}.$$

In homogeneous coordinates the projection matrix becomes 4 by 4 (but the origin doesn't move):

$$\text{Projection onto the plane } \mathbf{n}^T \mathbf{v} = 0 \quad P = \begin{bmatrix} I - \mathbf{n}\mathbf{n}^T & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}.$$

Now project onto a plane $\mathbf{n}^T(\mathbf{v} - \mathbf{v}_0) = 0$ that does *not* go through the origin. One point on the plane is \mathbf{v}_0 . This is an affine space (or a *flat*). It is like the solutions to $A\mathbf{v} = \mathbf{b}$ when the right side is not zero. One particular solution \mathbf{v}_0 is added to the nullspace—to produce a flat.

The projection onto the flat has three steps. Translate \mathbf{v}_0 to the origin by T_- . Project along the \mathbf{n} direction, and translate back along the row vector \mathbf{v}_0 :

$$\text{Projection onto a flat} \quad T_- P T_+ = \begin{bmatrix} I & 0 \\ -\mathbf{v}_0 & 1 \end{bmatrix} \begin{bmatrix} I - \mathbf{n}\mathbf{n}^T & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} I & 0 \\ \mathbf{v}_0 & 1 \end{bmatrix}.$$

I can't help noticing that T_- and T_+ are inverse matrices: translate and translate back. They are like the elementary matrices of Chapter 2.

The exercises will include reflection matrices, also known as *mirror matrices*. These are the fifth type needed in computer graphics. A reflection moves each point twice as far as a projection—the reflection goes through the plane and out the other side. So change the projection $I - \mathbf{n}\mathbf{n}^T$ to $I - 2\mathbf{n}\mathbf{n}^T$ for a mirror matrix.

The matrix P gave a “parallel” projection. All points move parallel to \mathbf{n} , until they reach the plane. The other choice in computer graphics is a “perspective” projection. This is more popular because it includes foreshortening. With perspective, an object looks larger as it moves closer. Instead of staying parallel to \mathbf{n} (and parallel to each other), the lines of projection come *toward the eye*—the center of projection. This is how we perceive depth in a two-dimensional photograph.

The basic problem of computer graphics starts with a scene and a viewing position. Ideally, the image on the screen is what the viewer would see. The simplest image assigns just one bit to every small picture element—called a *pixel*. It is light or dark. This gives a black and white picture with no shading. You would not approve. In practice, we assign shading levels between 0 and 2^8 for three colors like red, green, and blue. That means $8 \times 3 = 24$ bits for each pixel. Multiply by the number of pixels, and a lot of memory is needed!

Physically, a *raster frame buffer* directs the electron beam. It scans like a television set. The quality is controlled by the number of pixels and the number of bits per pixel. In this area, the standard text is *Computer Graphics: Principles and Practice* by Hughes, Van Dam, McGuire, Skylar, Foley, Feiner, and Akeley (3rd edition, Addison-Wesley, 2014). Notes by Ronald Goldman and by Tony DeRose were excellent references.

■ REVIEW OF THE KEY IDEAS ■

1. Computer graphics needs shift operations $T(\mathbf{v}) = \mathbf{v} + \mathbf{v}_0$ as well as linear operations $T(\mathbf{v}) = A\mathbf{v}$.
2. A shift in \mathbf{R}^n can be executed by a matrix of order $n + 1$, using homogeneous coordinates.
3. The extra component 1 in $[x \ y \ z \ 1]$ is preserved when all matrices have the numbers 0, 0, 0, 1 as last column.

Problem Set 10.6

- 1 A typical point in \mathbf{R}^3 is $x\mathbf{i} + y\mathbf{j} + z\mathbf{k}$. The coordinate vectors \mathbf{i} , \mathbf{j} , and \mathbf{k} are $(1, 0, 0)$, $(0, 1, 0)$, $(0, 0, 1)$. The coordinates of the point are (x, y, z) .
This point in computer graphics is $x\mathbf{i} + y\mathbf{j} + z\mathbf{k} + \mathbf{origin}$. Its homogeneous coordinates are $(\quad, \quad, \quad, \quad)$. Other coordinates for the same point are $(\quad, \quad, \quad, \quad)$.
- 2 A linear transformation T is determined when we know $T(\mathbf{i})$, $T(\mathbf{j})$, $T(\mathbf{k})$. For an affine transformation we also need $T(\quad)$. The input point $(x, y, z, 1)$ is transformed to $xT(\mathbf{i}) + yT(\mathbf{j}) + zT(\mathbf{k}) + \quad$.
- 3 Multiply the 4 by 4 matrix T for translation along $(1, 4, 3)$ and the matrix T_1 for translation along $(0, 2, 5)$. The product TT_1 is translation along \quad .
- 4 Write down the 4 by 4 matrix S that scales by a constant c . Multiply ST and also TS , where T is translation by $(1, 4, 3)$. To blow up the picture around the center point $(1, 4, 3)$, would you use $\mathbf{v}ST$ or $\mathbf{v}TS$?
- 5 What scaling matrix S (in homogeneous coordinates, so 3 by 3) would produce a 1 by 1 square page from a standard 8.5 by 11 page?
- 6 What 4 by 4 matrix would move a corner of a cube to the origin and then multiply all lengths by 2? The corner of the cube is originally at $(1, 1, 2)$.
- 7 When the three matrices in equation 1 multiply the unit vector \mathbf{a} , show that they give $(\cos \theta)\mathbf{a}$ and $(1 - \cos \theta)\mathbf{a}$ and $\mathbf{0}$. Addition gives $\mathbf{a}Q = \mathbf{a}$ and the rotation axis is not moved.
- 8 If \mathbf{b} is perpendicular to \mathbf{a} , multiply by the three matrices in 1 to get $(\cos \theta)\mathbf{b}$ and $\mathbf{0}$ and a vector perpendicular to \mathbf{b} . So $Q\mathbf{b}$ makes an angle θ with \mathbf{b} . **This is rotation.**
- 9 What is the 3 by 3 projection matrix $I - \mathbf{nn}^T$ onto the plane $\frac{2}{3}x + \frac{2}{3}y + \frac{1}{3}z = 0$? In homogeneous coordinates add 0, 0, 0, 1 as an extra row and column in P .

- 10 With the same 4 by 4 matrix P , multiply T_-PT_+ to find the projection matrix onto the plane $\frac{2}{3}x + \frac{2}{3}y + \frac{1}{3}z = 1$. The translation T_- moves a point on that plane (choose one) to $(0, 0, 0, 1)$. The inverse matrix T_+ moves it back.
- 11 Project $(3, 3, 3)$ onto those planes. Use P in Problem 9 and T_-PT_+ in Problem 10.
- 12 If you project a square onto a plane, what shape do you get?
- 13 If you project a cube onto a plane, what is the outline of the projection? Make the projection plane perpendicular to a diagonal of the cube.
- 14 The 3 by 3 mirror matrix that reflects through the plane $\mathbf{n}^T \mathbf{v} = 0$ is $M = I - 2\mathbf{n}\mathbf{n}^T$. Find the reflection of the point $(3, 3, 3)$ in the plane $\frac{2}{3}x + \frac{2}{3}y + \frac{1}{3}z = 0$.
- 15 Find the reflection of $(3, 3, 3)$ in the plane $\frac{2}{3}x + \frac{2}{3}y + \frac{1}{3}z = 1$. Take three steps T_-MT_+ using 4 by 4 matrices: translate by T_- so the plane goes through the origin, reflect the translated point $(3, 3, 3, 1)T_-$ in that plane, then translate back by T_+ .
- 16 The vector between the origin $(0, 0, 0, 1)$ and the point $(x, y, z, 1)$ is the difference $\mathbf{v} = \underline{\hspace{1cm}}$. In homogeneous coordinates, vectors end in $\underline{\hspace{1cm}}$. So we add a $\underline{\hspace{1cm}}$ to a point, not a point to a point.
- 17 If you multiply only the *last* coordinate of each point to get (x, y, z, c) , you rescale the whole space by the number $\underline{\hspace{1cm}}$. This is because the point (x, y, z, c) is the same as $(\underline{\hspace{1cm}}, \underline{\hspace{1cm}}, \underline{\hspace{1cm}}, 1)$.